

Bidirectional Dijkstra

May 30, 2020

Dijkstra's algorithm computes lengths of shortest paths from a start vertex s to every other vertex in a weighted graph with nonnegative weights. It works by successively improving an approximation $d[v]$ to the shortest path length $\delta(s, v)$ from s to v , which is initially $d[s] = 0$ and $d[v] = \infty$ for $v \neq s$. The algorithm maintains a priority queue Q of vertices which haven't yet been processed, initially containing just s with priority 0, and a set S of vertices whose true distance to s is known. It works as follows, in pseudo-Python:

```
while Q is not empty:
    u = extract_min(Q)
    S.add(u)
    for v in adj(u):
        if v not in S and d[v] > d[u] + w(u, v):
            d[v] = d[u] + w(u, v)
            Q.add(v, priority=d[v])
```

...where $\text{adj}(v)$ is the vertices adjacent to v and $w(u, v)$ is the weight of the edge $u \rightarrow v$. The function `extract_min` pops a vertex with minimum d -value out of Q . The procedure inside the conditional is called *relaxing* the edge (u, v) .

The key result in proving correctness is that once a vertex u enters S , its d -value $d[u]$ is equal to the true shortest path distance $\delta(s, u)$ - this is [CLRS Theorem 24.6](#).

Dijkstra runs in $O(V \log V + E \log V)$ if Q is a minheap, which can be improved slightly with fancier data structures. If all we care about is $\delta(s, t)$ for some fixed t we can speed it up for certain graphs as follows. Suppose each vertex has outvalency about m , and the edge distance from s to t is n . Then if the edge weights are roughly equal we expect to relax on the order of m^n edges by the time we reach t . On the other hand, if we run two searches, one starting at s and the other at t , stopping when they meet in the middle, we would only relax $2m^{n/2}$ edges. This method is called **bidirectional Dijkstra**.

The subtlety in bidirectional Dijkstra is the stopping condition. When we detect an edge between the forward and backward sets S_f, S_b of processed vertices there is no guarantee a shortest path $s \rightsquigarrow t$ passes through that edge. "Stop as soon as you find an edge (u, v)

between vertices processed in the forward and backward searches, return $d_f[u] + w(u, v) + d_b[v]$ or $d_f[v] + w(u, v) + d_b[u]$ as appropriate" is not correct.

(here d_f is the approximation to $\delta(s, -)$ computed by the forward search and d_b the approximation to $\delta(-, t)$ computed by the backward search).

The point of this post is to give a description of bidirectional Dijkstra that is precise enough to implement, then a sketch of why it is correct. This isn't in CLRS, for example, and while there are slides on the internet (e.g. [here](#) or [here](#) or [these 2016 recitation notes from 6.006](#)) the explanations have gaps which would have been filled in by the lecturer.

I'm going to assume that the graph is undirected; if it is directed you need to do the backward search on the opposite graph.

We start with forward approximations $d_f[v]$ of the distance of a node from s and backward approximations $d_b[v]$ of the distance of a node to t , initially all infinity except $d_f[s] = 0 = d_b[t]$. We keep a forward priority queue Q_f and a backward priority queue Q_b , initially containing s and t respectively, and sets S_f and S_b of vertices processed in the forward and backward searches, initially empty. The priority of an element in the queue is its d_f or d_b value. We also keep a number μ equal to the length of the shortest path $s \rightsquigarrow t$ yet seen, initially infinity.

```
while Qf is not empty and Qb is not empty:
    u = extract_min(Qf); v = extract_min(Qb)
    Sf.add(u); Sb.add(v)
    for x in adj(u):
        relax(u, x)
        if x in Sb and df[u] + w(u, x) + db[x] < mu:
            mu = df[u] + w(u, x) + db[x]
    for x in adj(v):
        relax(v, x)
        if x in Sf and db[v] + w(v, x) + df[x] < mu:
            mu = db[v] + w(v, x) + df[x]
    if df[u] + db[v] >= mu:
        break # mu is the true distance s-t
```

(Thanks to Panagiotis Karras, Ciaperoni Martino, Nassos Katsamanis, and Aristides Gionis who emailed me with a correction to this algorithm). You can find example Python code with tests [in my Github repository](#).

"Relax" refers to the same procedure as in the ordinary Dijkstra algorithm, so that for example `relax(u, x)` would do the following:

```

if (x is not in Sf) and df[x] > df[u] + weight(u, x):
    df[x] = df[u] + weight(u, x)
    Qf.add(x, priority=df[x])

```

I claim that when this algorithm leaves the while loop, μ equals the true distance $\delta(s, t)$. If you want to recover the actual shortest path you need to maintain a vertex as well as a best path length, and when μ gets updated also update that vertex to x . The shortest path is then a shortest path $s \rightsquigarrow x$ followed by a shortest path $x \rightsquigarrow t$.

First we need a result about *ordinary* Dijkstra.

At any stage of the ordinary Dijkstra algorithm, if $y \notin S$ then $\delta(y, s) \geq \max_{x \in S} \delta(x, s)$.

If not there is a first relaxation in the algorithm setting the d -value of a vertex $y \in Q$ to a value less than $d[x]$ for some $x \in S$. Say the edge (r, y) was the one being relaxed, so r was just added to S . $d[y]$ gets set to something larger than $\lambda = d[r]$, so $d[r] < d[y] < d[x]$.


If x was in S when $d[r]$ got set to λ , we have a contradiction to the relaxation above being the earliest setting a d -value of something in Q to something less than a d -value in S . So the sequence of events must have been:

1. $d[r]$ was set to λ
2. x entered S (before r did)
3. (r, y) was relaxed when r entered S . This is impossible: at step 2, x can't have had the smallest d -value outside S because $d[r]$ was already $\lambda < d[x]$.

Now we're ready to prove the correctness of the termination condition in bidirectional Dijkstra. Suppose that the algorithm terminated at the `break` but that there is a path $s \overset{P}{\rightsquigarrow} t$ beating μ . It can't contain a vertex x outside $S_f \cup S_b$: such a vertex is at least $d_f[u]$ from s and $d_b[v]$ from t by the ordinary Dijkstra result above, so P would have length at least μ . Thus P is contained in $S_f \cup S_b$ - but then it was considered when one of its edges was first found to connect the two sets of processed vertices, and μ was updated to something at most the length of P .

Matthew Towers' homepage

Matthew Towers' homepage
m.towers@ucl.ac.uk

 [matthew-towers](https://github.com/matthew-towers)

Hello world.

